

# YIFAN JIANG

yifanjiang97@utexas.edu | Website: [yifanjiang19.github.io](https://yifanjiang19.github.io) | Google Scholar Citation: 6869

## EDUCATION

**The University of Texas at Austin**, Austin, USA

2020 – 2024

*Ph.D.* in Electrical and Computer Engineering<sup>1</sup>

- Advisor: [Zhangyang \(Atlas\) Wang](#)

Research Interests: Machine Learning, Generative Model

**Huazhong University of Science and Technology**, Wuhan, China

2015 – 2019

*B.E.* in Electronic Information Engineering

## INDUSTRIAL EXPERIENCE

**Meta**, Menlo Park, USA

Jun. 2025 - Present

*Senior Research Scientist at Meta SuperIntelligence Lab*

Core contributor of the media generation model, delivered multiple key features to the [Meta AI/Vibes App](#).

- **Video Storytelling.** Advanced "Storytelling" feature in the Meta AI/Vibes from concept to the dogfooding stage.
  - Developed the training and inference infrastructure to support a **33B** text-to-videoaudio model. Implemented **FSDP / HSDP / CP** for robust training. Optionally supporting LoRA for efficiency.
  - Built data pipeline to collect and filter human-centric video data. Designed both high-level signals (VLM/LLM-based, face recognition, etc.) and low-level signals (aesthetic score, motion score, etc.).
  - Designed training recipes on both **Mid-training and SFT stages**, using human rating as signals for hyper-parameter tuning.
  - Constructed an LLM-based rewriter for prompt enhancement.
- **Person Insertion.** Led the development of the image editing feature "Add Me", successfully launching it in the Meta AI / Vibes app, beating Nano Banana Pro in both internal benchmarks and public A/B tests.
  - Architected a large-scale data filtering pipeline to support **Mid-training and SFT stages**, covering different distributions (synthetic and real).
  - Formulated training recipes for Mid-training and SFT stages, ablating different data mixture ratios. Orchestrated with distillation team to deliver a 12-step diffusion model.
  - Constructed an evaluation pipeline utilizing human annotators.
- **Personalized Photo Generation.** Adopted reinforcement learning to improve the quality of generated contents.
  - Developed a **GRPO-based reinforcement learning** pipeline for diffusion model.
  - Collected data and trained a visual quality assessment model as the **reward model**, to evaluate content quality.

**Apple**, Seattle, USA

Mar. 2024 – Jun. 2025

*Senior Research Scientist at AI/ML DMLI*

Core contributor of the Apple Video generative model and 3D generative model, responsible for training and modeling.

- **Text-to-Video Model.** Conducted ablation studies on different model recipes from 1B to 15B parameters. Curated SFT data for better motion and aesthetic.
- **3D Diffusion Model.** Architected a multi-view attention operator to improve 3D consistency. Optimized training speed utilizing PyTorch Profiler. Developed specialized noise scheduler for faster convergence.

**Apple**, Seattle, USA

May. 2023 – Jan. 2024

*Research Intern at AI/ML Argos Team, worked with [Liangliang Cao](#)*

- Designed an efficient diffusion model training algorithm for novel view synthesis applications.

**Adobe**, San Jose, USA

May. 2022 – Dec. 2022

*Research Intern at [Marc Levoy's Team](#)*

- Developed a high-quality depth estimator using multi-sensor assistance.

<sup>1</sup>Studied at Texas A&M University from Aug. 2019 to Aug. 2020; then transferred with my advisor to UT Austin

**Google Research**, Mountain View, USA

May. 2021 – May. 2022

Research Intern at *GCam*, mentored by *Tianfan Xue*, *Ben Mildenhall*, and *Jon Barron*

- Developed a fast denoising operator and an efficient backbone.
- Designed a high-fidelity neural radiance field that improves the rendering quality.

**Adobe**, San Jose, USA

May. 2020 – Nov. 2020

Research Intern at *Applied Research Team (ART)*

- Developed a self-supervised method for image harmonization.

**Bytedance AI Lab**, Beijing, China

Jan. 2019 – Aug. 2019

Research Intern at *Bytedance AI Lab*, mentored by *Xiaohui Shen* and *Jianchao Yang*

- Designed a jointly image denoising and low-light enhancement algorithm.

## SELECTED PUBLICATIONS

---

(\* indicates equal contribution)

- [ICLR] **Y. Jiang**, H. Tang, J. Chang, D. Xu, L. Song, Z. Wang and L. Cao, “Efficient-3DiM: Learning a Generalizable Single-image Novel-view Synthesizer in One Day”, *International Conference on Learning Representations*, 2024
- [CVPR] D. Xu, **Y. Jiang**, P. Wang, Z. Fan, Y. Wang, and Z. Wang, “NeuralLift-360: Lifting An In-the-wild 2D Photo to A 3D Object with 360° Views”, *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2023. **Highlight (2.5% of 9155 submissions)**
- [CVPR] **Y. Jiang**, H. Peter, B. Mildenhall, D. Xu, J. T. Barron, and Z. Wang, “AlignNeRF: High-Fidelity Neural Radiance Fields via Alignment-Aware Training”, *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2023
- [ECCV] **Y. Jiang\***, D. Xu\*, P. Wang, Z. Fan, H. Shi, and Z. Wang. “SinNeRF: Training Neural Radiance Fields on Complex Scenes from a Single Image”, *European Conference on Computer Vision*, 2022.
- [ECCV] **Y. Jiang\***, Z. Fan\*, P. Wang\*, X. Gong, D. Xu, and Z. Wang. “Unified Implicit Neural Stylization”, *European Conference on Computer Vision*, 2022.
- [ECCV] **Y. Jiang**, B. Wronski, B. Mildenhall, J. T. Barron, Z. Wang, and T. Xue. “Fast and High-Quality Image Denoising via Malleable Convolutions”, *European Conference on Computer Vision*, 2022.
- [NeurIPS] **Y. Jiang**, S. Chang, and Z. Wang, “TransGAN: Two Pure Transformers can Make One Strong GAN and That Can Scale Up”, *Advances in Neural Information Processing Systems*, 2021.
- [TIP] **Y. Jiang**, X. Gong, D. Liu, Y. Cheng, C. Fang, X. Shen, J. Yang, P. Zhou, and Z. Wang, “EnlightenGAN: Deep Light Enhancement without Paired Supervision”, *IEEE Transaction on Image Processing*. **IEEE SPS Young Author Best Paper Award.**

## MEDIA HIGHLIGHT

---

- TransGAN was covered by *Quanta Magazine* (Mar. 2022) and was highlighted by *Top AI influencers* and high-profile *YouTubers*, as well as considered as *the most influential new paper of the month (Feb. 2021)*.
- AutoGAN was covered by *Synced AI Technology & Industry Review* (Aug. 2019), and also featured on *Towards Data Science* (Sep. 2019) and *Analytics Magazine* (Aug. 2019), etc.

## COMMUNITY SERVICES

---

- Reviewer for: CVPR, ICCV, ECCV, ICML, NeurIPS, ICLR, Siggraph, Siggraph Aisa, IROS, WACV, IJCAI, TPAMI, TIP, IJCV, NeuroComputing, RA-L, TCI, TCSVT
- Workshop Organizer for: *ECCV RLQ-TOD Workshop 2020*

## AWARDS

---

- **IEEE SPS Young Author Best Paper Award** 2024
- **University of Texas Graduate Dean’s Prestigious Fellowship** 2023
- **Apple Scholars in AI/ML Fellowship** 2023